

CV_template

Erie Seong Ho Han

12/6/2021

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.3.1 --
```

```
## v ggplot2 3.3.3    v purrr  0.3.4
## v tibble  3.1.4    v dplyr  1.0.7
## v tidyr   1.1.3    v stringr 1.4.0
## v readr   2.0.1    v forcats 0.5.1
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

```
library(glmnet)
```

```
## Loading required package: Matrix
```

```
##
```

```
## Attaching package: 'Matrix'
```

```
## The following objects are masked from 'package:tidyr':
```

```
##
```

```
##     expand, pack, unpack
```

```
## Loaded glmnet 4.1-2
```

```
library(stringr)
```

```
library(bnstruct)
```

```
## Loading required package: bitops
```

```
##
```

```
## Attaching package: 'bitops'
```

```
## The following object is masked from 'package:Matrix':
```

```
##
```

```
##     %&%
```

```

## Loading required package: igraph

##
## Attaching package: 'igraph'

## The following objects are masked from 'package:dplyr':
##
##   as_data_frame, groups, union

## The following objects are masked from 'package:purrr':
##
##   compose, simplify

## The following object is masked from 'package:tidyr':
##
##   crossing

## The following object is masked from 'package:tibble':
##
##   as_data_frame

## The following objects are masked from 'package:stats':
##
##   decompose, spectrum

## The following object is masked from 'package:base':
##
##   union

##
## Attaching package: 'bnstruct'

## The following object is masked from 'package:tidyr':
##
##   complete

library(randomForest)

## randomForest 4.6-14

## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'

## The following object is masked from 'package:dplyr':
##
##   combine

## The following object is masked from 'package:ggplot2':
##
##   margin

```

```
library(e1071)
```

```
##  
## Attaching package: 'e1071'  
  
## The following object is masked from 'package:bnstruct':  
##  
##      impute
```

Including Plots

```
heart <- read.csv(file = 'heart.csv')  
heart %>%  
  count(Cholesterol)
```

```
##      Cholesterol    n  
## 1              0 172  
## 2             85    1  
## 3            100    2  
## 4            110    1  
## 5            113    1  
## 6            117    1  
## 7            123    1  
## 8            126    2  
## 9            129    1  
## 10           131    1  
## 11           132    1  
## 12           139    2  
## 13           141    1  
## 14           142    1  
## 15           147    2  
## 16           149    2  
## 17           152    1  
## 18           153    1  
## 19           156    1  
## 20           157    1  
## 21           159    1  
## 22           160    6  
## 23           161    2  
## 24           163    2  
## 25           164    2  
## 26           165    1  
## 27           166    4  
## 28           167    3  
## 29           168    2  
## 30           169    2  
## 31           170    2  
## 32           171    3  
## 33           172    2  
## 34           173    2
```

## 35	174	1
## 36	175	4
## 37	176	1
## 38	177	6
## 39	178	1
## 40	179	2
## 41	180	3
## 42	181	2
## 43	182	5
## 44	183	1
## 45	184	4
## 46	185	3
## 47	186	6
## 48	187	2
## 49	188	4
## 50	190	2
## 51	192	4
## 52	193	6
## 53	194	2
## 54	195	7
## 55	196	6
## 56	197	7
## 57	198	6
## 58	199	3
## 59	200	4
## 60	201	6
## 61	202	3
## 62	203	7
## 63	204	9
## 64	205	3
## 65	206	3
## 66	207	6
## 67	208	7
## 68	209	5
## 69	210	4
## 70	211	9
## 71	212	6
## 72	213	7
## 73	214	7
## 74	215	6
## 75	216	9
## 76	217	4
## 77	218	6
## 78	219	8
## 79	220	10
## 80	221	5
## 81	222	6
## 82	223	10
## 83	224	6
## 84	225	7
## 85	226	6
## 86	227	4
## 87	228	5
## 88	229	4

## 89	230	9
## 90	231	5
## 91	232	3
## 92	233	6
## 93	234	7
## 94	235	5
## 95	236	6
## 96	237	6
## 97	238	4
## 98	239	4
## 99	240	8
## 100	241	4
## 101	242	2
## 102	243	7
## 103	244	4
## 104	245	6
## 105	246	8
## 106	247	3
## 107	248	6
## 108	249	5
## 109	250	5
## 110	251	1
## 111	252	3
## 112	253	4
## 113	254	11
## 114	255	3
## 115	256	5
## 116	257	3
## 117	258	7
## 118	259	2
## 119	260	8
## 120	261	3
## 121	262	1
## 122	263	8
## 123	264	6
## 124	265	4
## 125	266	4
## 126	267	5
## 127	268	5
## 128	269	6
## 129	270	6
## 130	271	4
## 131	272	3
## 132	273	5
## 133	274	6
## 134	275	7
## 135	276	4
## 136	277	5
## 137	278	1
## 138	279	1
## 139	280	2
## 140	281	3
## 141	282	7
## 142	283	5

## 143	284	4
## 144	285	2
## 145	286	2
## 146	287	2
## 147	288	6
## 148	289	6
## 149	290	2
## 150	291	3
## 151	292	4
## 152	293	1
## 153	294	4
## 154	295	5
## 155	297	4
## 156	298	5
## 157	299	2
## 158	300	2
## 159	302	2
## 160	303	4
## 161	304	2
## 162	305	4
## 163	306	3
## 164	307	2
## 165	308	6
## 166	309	4
## 167	310	3
## 168	311	2
## 169	312	2
## 170	313	1
## 171	315	3
## 172	316	1
## 173	318	3
## 174	319	1
## 175	320	2
## 176	321	1
## 177	322	1
## 178	325	2
## 179	326	2
## 180	327	1
## 181	328	1
## 182	329	1
## 183	330	2
## 184	331	1
## 185	333	1
## 186	335	2
## 187	336	1
## 188	337	1
## 189	338	1
## 190	339	2
## 191	340	2
## 192	341	3
## 193	342	3
## 194	344	1
## 195	347	1
## 196	349	1

## 197	353	1
## 198	354	1
## 199	355	1
## 200	358	1
## 201	360	1
## 202	365	1
## 203	369	1
## 204	384	1
## 205	385	1
## 206	388	1
## 207	392	1
## 208	393	1
## 209	394	2
## 210	404	1
## 211	407	1
## 212	409	1
## 213	412	1
## 214	417	1
## 215	458	1
## 216	466	1
## 217	468	1
## 218	491	1
## 219	518	1
## 220	529	1
## 221	564	1
## 222	603	1

```
heart %>%
  count(RestingBP)
```

##	RestingBP	n
## 1	0	1
## 2	80	1
## 3	92	1
## 4	94	2
## 5	95	6
## 6	96	1
## 7	98	1
## 8	100	15
## 9	101	1
## 10	102	3
## 11	104	3
## 12	105	9
## 13	106	3
## 14	108	7
## 15	110	58
## 16	112	14
## 17	113	1
## 18	114	2
## 19	115	19
## 20	116	2
## 21	117	1
## 22	118	10
## 23	120	132

```
## 24      122  12
## 25      123   2
## 26      124  12
## 27      125  29
## 28      126   7
## 29      127   1
## 30      128  18
## 31      129   1
## 32      130 118
## 33      131   4
## 34      132  17
## 35      133   6
## 36      134  11
## 37      135  20
## 38      136  13
## 39      137   5
## 40      138  17
## 41      139   5
## 42      140 107
## 43      141   3
## 44      142  11
## 45      143   2
## 46      144   8
## 47      145  18
## 48      146   4
## 49      148   2
## 50      150  55
## 51      152   7
## 52      154   3
## 53      155   8
## 54      156   2
## 55      158   4
## 56      160  50
## 57      164   1
## 58      165   2
## 59      170  14
## 60      172   2
## 61      174   1
## 62      178   3
## 63      180  12
## 64      185   1
## 65      190   2
## 66      192   1
## 67      200   4
```

```
heart %>%
  count(Age)
```

```
##   Age  n
## 1  28  1
## 2  29  3
## 3  30  1
## 4  31  2
## 5  32  5
```



```
## 6 33 2
## 7 34 7
## 8 35 11
## 9 36 6
## 10 37 11
## 11 38 16
## 12 39 15
## 13 40 13
## 14 41 24
## 15 42 18
## 16 43 24
## 17 44 19
## 18 45 18
## 19 46 24
## 20 47 19
## 21 48 31
## 22 49 21
## 23 50 25
## 24 51 35
## 25 52 36
## 26 53 33
## 27 54 51
## 28 55 41
## 29 56 38
## 30 57 38
## 31 58 42
## 32 59 35
## 33 60 32
## 34 61 31
## 35 62 35
## 36 63 30
## 37 64 22
## 38 65 21
## 39 66 13
## 40 67 15
## 41 68 10
## 42 69 13
## 43 70 7
## 44 71 5
## 45 72 4
## 46 73 1
## 47 74 7
## 48 75 3
## 49 76 2
## 50 77 2
```

```
heart %>%
  count(MaxHR)
```

```
##      MaxHR  n
## 1       60  1
## 2       63  1
## 3       67  1
## 4       69  1
```

## 5	70	1
## 6	71	1
## 7	72	2
## 8	73	1
## 9	77	1
## 10	78	1
## 11	80	2
## 12	82	3
## 13	83	1
## 14	84	3
## 15	86	4
## 16	87	1
## 17	88	2
## 18	90	3
## 19	91	1
## 20	92	6
## 21	93	2
## 22	94	4
## 23	95	2
## 24	96	7
## 25	97	3
## 26	98	9
## 27	99	7
## 28	100	14
## 29	102	4
## 30	103	4
## 31	104	2
## 32	105	11
## 33	106	5
## 34	107	1
## 35	108	8
## 36	109	5
## 37	110	23
## 38	111	5
## 39	112	13
## 40	113	5
## 41	114	6
## 42	115	16
## 43	116	9
## 44	117	6
## 45	118	12
## 46	119	5
## 47	120	36
## 48	121	5
## 49	122	20
## 50	123	7
## 51	124	9
## 52	125	21
## 53	126	12
## 54	127	8
## 55	128	14
## 56	129	4
## 57	130	33
## 58	131	7

## 59	132	11
## 60	133	5
## 61	134	6
## 62	135	15
## 63	136	6
## 64	137	7
## 65	138	14
## 66	139	6
## 67	140	41
## 68	141	6
## 69	142	14
## 70	143	10
## 71	144	13
## 72	145	14
## 73	146	6
## 74	147	5
## 75	148	11
## 76	149	6
## 77	150	43
## 78	151	5
## 79	152	11
## 80	153	5
## 81	154	12
## 82	155	14
## 83	156	10
## 84	157	7
## 85	158	8
## 86	159	5
## 87	160	25
## 88	161	7
## 89	162	13
## 90	163	10
## 91	164	4
## 92	165	11
## 93	166	5
## 94	167	2
## 95	168	8
## 96	169	6
## 97	170	20
## 98	171	4
## 99	172	10
## 100	173	7
## 101	174	7
## 102	175	10
## 103	176	2
## 104	177	1
## 105	178	6
## 106	179	6
## 107	180	10
## 108	181	2
## 109	182	6
## 110	184	4
## 111	185	4
## 112	186	2

```
## 113 187 1
## 114 188 2
## 115 190 2
## 116 192 1
## 117 194 1
## 118 195 1
## 119 202 1
```

```
heart %>%
  count(Oldpeak)
```

```
##   Oldpeak    n
## 1   -2.6     1
## 2   -2.0     1
## 3   -1.5     1
## 4   -1.1     1
## 5   -1.0     2
## 6   -0.9     1
## 7   -0.8     1
## 8   -0.7     1
## 9   -0.5     2
## 10  -0.1     2
## 11    0.0  368
## 12    0.1    14
## 13    0.2    22
## 14    0.3    11
## 15    0.4    11
## 16    0.5    19
## 17    0.6    14
## 18    0.7     7
## 19    0.8    16
## 20    0.9     4
## 21    1.0    86
## 22    1.1     7
## 23    1.2    26
## 24    1.3     7
## 25    1.4    18
## 26    1.5    53
## 27    1.6    16
## 28    1.7     6
## 29    1.8    17
## 30    1.9     7
## 31    2.0    76
## 32    2.1     2
## 33    2.2     5
## 34    2.3     2
## 35    2.4     4
## 36    2.5    16
## 37    2.6     7
## 38    2.8     7
## 39    2.9     1
## 40    3.0    28
## 41    3.1     1
## 42    3.2     2
```

```
## 43      3.4    3
## 44      3.5    2
## 45      3.6    4
## 46      3.7    1
## 47      3.8    1
## 48      4.0    8
## 49      4.2    2
## 50      4.4    1
## 51      5.0    1
## 52      5.6    1
## 53      6.2    1
```

```
heart %>%
  count(ST_Slope)
```

```
##   ST_Slope    n
## 1     Down   63
## 2     Flat 460
## 3      Up 395
```

```
heart %>%
  count(HeartDisease)
```

```
##   HeartDisease    n
## 1             0 410
## 2             1 508
```

```
heart <- heart %>%
  filter(Cholesterol != 0) %>%
  filter(RestingBP!=0) %>%
  mutate(Sex = as.factor(Sex)) %>%
  mutate(ChestPainType= as.factor(ChestPainType)) %>%
  mutate(RestingECG= as.factor(RestingECG)) %>%
  mutate(ExerciseAngina = as.factor(ExerciseAngina)) %>%
  mutate(ST_Slope = as.factor(ST_Slope)) %>%
  mutate(HeartDisease = as.factor(HeartDisease))
```

```
set.seed(1)
shuffled_heart <- heart[sample(nrow(heart)),]
folds <- cut(seq(1,nrow(shuffled_heart)),breaks=5,labels=FALSE)
```

```
# error
misclassification.linear <- rep(0, 5)
misclassification.poly <- rep(0, 5)
misclassification.radial <- rep(0, 5)

get_misclassification <- function(bestmod, X_test, y_test) {
  prediction <- predict(bestmod, X_test)
  tab <- table(y_test, prediction)
  misclassification_rate <- 1 - sum(diag(tab))/sum(tab)
  return(misclassification_rate)
}
```

```

# Cross validation
for(i in 1:5){

  #Segment your data by fold using the which() function
  testIndexes <- which(folds==i,arr.ind=TRUE)
  testData <- shuffled_heart[testIndexes, ]
  y.test <- testData$HeartDisease
  X.test <- testData[, -12]
  trainData <- shuffled_heart[-testIndexes, ]

  # Need to choose parameters
  set.seed(1)
  tune.out.linear <- tune(svm, HeartDisease ~ ., data = trainData, kernel = "linear",
    ranges = list(cost = 10^seq(-2, 1, by = 0.25)))
  tune.out.poly <- tune(svm, HeartDisease ~ ., data = trainData, kernel = "poly", degree=2,
    ranges = list(cost = 10^seq(-2, 1, by = 0.25)))
  tune.out.radial <- tune(svm, HeartDisease ~ ., data = trainData, kernel = "radial",
    ranges = list(cost = 10^seq(-2, 1, by = 0.25)))

  bestmod.linear <- tune.out.linear$best.model
  bestmod.poly <- tune.out.poly$best.model
  bestmod.radial <- tune.out.radial$best.model

  misclassification.linear[i] = get_misclassification(bestmod.linear, X.test, y.test)
  misclassification.poly[i] = get_misclassification(bestmod.poly, X.test, y.test)
  misclassification.radial[i] = get_misclassification(bestmod.radial, X.test, y.test)
}

misclassification.linear

## [1] 0.1400000 0.1208054 0.1610738 0.1140940 0.1275168

misclassification.poly

## [1] 0.1666667 0.1342282 0.1744966 0.1275168 0.1073826

misclassification.radial

## [1] 0.1533333 0.1275168 0.1812081 0.1476510 0.1208054

mean(misclassification.linear)

## [1] 0.132698

mean(misclassification.poly)

## [1] 0.1420582

```

```
mean(misclassification.radial)
```

```
## [1] 0.1461029
```

linear SVM has the best estimated test accuracy on average.

```
# Tuning linear SVM
tune.out.linear <- tune(svm, HeartDisease ~ ., data = heart, kernel = "linear",
                        ranges = list(cost = 10^seq(-2, 1, by = 0.25)))
bestmod.linear <- tune.out.linear$best.model
bestmod.param <- tune.out.linear$best.parameters
bestmod.linear
```

```
##
## Call:
## best.tune(method = svm, train.x = HeartDisease ~ ., data = heart,
##   ranges = list(cost = 10^seq(-2, 1, by = 0.25)), kernel = "linear")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##       cost:  0.3162278
##
## Number of Support Vectors:  257
```

```
bestmod.param
```

```
##       cost
## 7 0.3162278
```

Best hyperparameter is $0.3152278 = 10^{-0.5}$

CV comparing all

```
# error
misclassification.svm <- rep(0, 5)
misclassification.rf <- rep(0, 5)
misclassification.glm <- rep(0, 5)
misclassification.final <- rep(0, 5)
misclassification.glm.aic <- rep(0,5)

get_misclassification_with_prediction <- function(prediction, y_test) {
  tab <- table(y_test, prediction)
  misclassification_rate <- 1 - sum(diag(tab))/sum(tab)
  return(misclassification_rate)
}
```

```

# Cross validation
for(i in 1:5){

  #Segment your data by fold using the which() function
  testIndexes <- which(folds==i,arr.ind=TRUE)
  testData <- shuffled_heart[testIndexes, ]
  y.test <- testData$HeartDisease
  X.test <- testData[, -12]
  trainData <- shuffled_heart[-testIndexes, ]

  # Need to choose parameters
  set.seed(1)
  svm.linear <- svm(HeartDisease ~ ., kernel = "linear", data = trainData, cost = 0.3152278)
  rf.fit <- randomForest(HeartDisease ~ ., data = trainData, mtry = sqrt(11), ntree = 1000)
  glm.fit <- glm(HeartDisease ~ Sex + Age + ChestPainType + RestingBP + Oldpeak + FastingBS+ ST_Slope+
    data = trainData)

  # From AIC
  glm.fit2 <- glm(formula = HeartDisease ~ Age + Sex + ChestPainType + RestingBP +
    Cholesterol + FastingBS + RestingECG + MaxHR + ExerciseAngina +
    Oldpeak + ST_Slope + Age:ST_Slope + Sex:FastingBS + Sex:MaxHR +
    Sex:ExerciseAngina + ChestPainType:Cholesterol + ChestPainType:FastingBS +
    ChestPainType:RestingECG + ChestPainType:ST_Slope + RestingBP:Oldpeak +
    Cholesterol:ST_Slope + FastingBS:ST_Slope + RestingECG:ExerciseAngina +
    MaxHR:ST_Slope, family = "binomial", data = trainData)

  prediction.svm <- predict(svm.linear, X.test)
  prediction.rf <- predict(rf.fit, X.test)

  # prediction with glm returns prediction for the logit, thus it seems that it's necessary
  # to convert that to the labels we need manually
  prediction.glm <- predict(glm.fit, X.test, type="response")
  prediction.glm2 <- predict(glm.fit2, X.test, type="response")
  for(j in 1:length(prediction.glm)) {
    if(prediction.glm[j] < 0.5) {
      prediction.glm[j] = 0
    }
    else {
      prediction.glm[j] = 1
    }

    if(prediction.glm2[j] < 0.5) {
      prediction.glm2[j] = 0
    }
    else {
      prediction.glm2[j] = 1
    }
  }

  prediction.total <- prediction.glm

```



```

#Compute the final model
for(j in 1:length(prediction.total)) {
  if(prediction.svm[j] == 1) {
    prediction.total[j] = 1
  }
  if(prediction.rf[j] == 1) {
    prediction.total[j] = prediction.total[j] + 1
  }
  if(prediction.glm2[j] == 1) {
    prediction.total[j] = prediction.total[j] + 1
  }
  if(prediction.total[j] > 1.5) {
    prediction.total[j] = 1
  }
}

misclassification.svm[i] = get_misclassification_with_prediction(prediction.svm, y.test)
misclassification.rf[i] = get_misclassification_with_prediction(prediction.rf, y.test)
misclassification.glm[i] = get_misclassification_with_prediction(prediction.glm, y.test)
misclassification.final[i] = get_misclassification_with_prediction(prediction.total, y.test)
misclassification.glm.aic[i] = get_misclassification_with_prediction(prediction.glm2, y.test)
}

```

```

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

```

```
misclassification.svm
```

```
## [1] 0.1400000 0.1208054 0.1812081 0.1073826 0.1208054
```

```
misclassification.rf
```

```
## [1] 0.1266667 0.1342282 0.1610738 0.1275168 0.1140940
```

```
misclassification.glm
```

```
## [1] 0.1933333 0.1476510 0.2147651 0.1610738 0.1140940
```

```
misclassification.final
```

```
## [1] 0.1733333 0.1342282 0.1812081 0.1073826 0.1476510
```

```
misclassification.glm.aic
```

```
## [1] 0.1466667 0.1476510 0.2147651 0.1342282 0.1073826
```

```
mean(misclassification.svm)
```

```
## [1] 0.1340403
```

```
mean(misclassification.rf)
```

```
## [1] 0.1327159
```

```
mean(misclassification.glm)
```

```
## [1] 0.1661834
```

```
mean(misclassification.final)
```

```
## [1] 0.1487606
```

```
mean(misclassification.glm.aic)
```

```
## [1] 0.1501387
```