

SVM

```
library(tidyverse)

## Warning in system("timedatectl", intern = TRUE): running command 'timedatectl'
## had status 1

## -- Attaching packages ----- tidyverse 1.3.1 --

## v ggplot2 3.3.5    v purrr  0.3.4
## v tibble  3.1.5    v dplyr  1.0.7
## v tidyr   1.1.4    v stringr 1.4.0
## v readr   2.0.2    v forcats 0.5.1

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()

library(stringr)
library(e1071)

heart <- read.csv(file = 'heart.csv')
heart %>%
  count(Cholesterol)

##      Cholesterol    n
## 1             0 172
## 2             85   1
## 3            100   2
## 4            110   1
## 5            113   1
## 6            117   1
## 7            123   1
## 8            126   2
## 9            129   1
## 10           131   1
## 11           132   1
## 12           139   2
## 13           141   1
## 14           142   1
## 15           147   2
## 16           149   2
## 17           152   1
## 18           153   1
## 19           156   1
## 20           157   1
## 21           159   1
## 22           160   6
## 23           161   2
## 24           163   2
## 25           164   2
```

## 26	165	1
## 27	166	4
## 28	167	3
## 29	168	2
## 30	169	2
## 31	170	2
## 32	171	3
## 33	172	2
## 34	173	2
## 35	174	1
## 36	175	4
## 37	176	1
## 38	177	6
## 39	178	1
## 40	179	2
## 41	180	3
## 42	181	2
## 43	182	5
## 44	183	1
## 45	184	4
## 46	185	3
## 47	186	6
## 48	187	2
## 49	188	4
## 50	190	2
## 51	192	4
## 52	193	6
## 53	194	2
## 54	195	7
## 55	196	6
## 56	197	7
## 57	198	6
## 58	199	3
## 59	200	4
## 60	201	6
## 61	202	3
## 62	203	7
## 63	204	9
## 64	205	3
## 65	206	3
## 66	207	6
## 67	208	7
## 68	209	5
## 69	210	4
## 70	211	9
## 71	212	6
## 72	213	7
## 73	214	7
## 74	215	6
## 75	216	9
## 76	217	4
## 77	218	6
## 78	219	8
## 79	220	10

## 80	221	5
## 81	222	6
## 82	223	10
## 83	224	6
## 84	225	7
## 85	226	6
## 86	227	4
## 87	228	5
## 88	229	4
## 89	230	9
## 90	231	5
## 91	232	3
## 92	233	6
## 93	234	7
## 94	235	5
## 95	236	6
## 96	237	6
## 97	238	4
## 98	239	4
## 99	240	8
## 100	241	4
## 101	242	2
## 102	243	7
## 103	244	4
## 104	245	6
## 105	246	8
## 106	247	3
## 107	248	6
## 108	249	5
## 109	250	5
## 110	251	1
## 111	252	3
## 112	253	4
## 113	254	11
## 114	255	3
## 115	256	5
## 116	257	3
## 117	258	7
## 118	259	2
## 119	260	8
## 120	261	3
## 121	262	1
## 122	263	8
## 123	264	6
## 124	265	4
## 125	266	4
## 126	267	5
## 127	268	5
## 128	269	6
## 129	270	6
## 130	271	4
## 131	272	3
## 132	273	5
## 133	274	6

## 134	275	7
## 135	276	4
## 136	277	5
## 137	278	1
## 138	279	1
## 139	280	2
## 140	281	3
## 141	282	7
## 142	283	5
## 143	284	4
## 144	285	2
## 145	286	2
## 146	287	2
## 147	288	6
## 148	289	6
## 149	290	2
## 150	291	3
## 151	292	4
## 152	293	1
## 153	294	4
## 154	295	5
## 155	297	4
## 156	298	5
## 157	299	2
## 158	300	2
## 159	302	2
## 160	303	4
## 161	304	2
## 162	305	4
## 163	306	3
## 164	307	2
## 165	308	6
## 166	309	4
## 167	310	3
## 168	311	2
## 169	312	2
## 170	313	1
## 171	315	3
## 172	316	1
## 173	318	3
## 174	319	1
## 175	320	2
## 176	321	1
## 177	322	1
## 178	325	2
## 179	326	2
## 180	327	1
## 181	328	1
## 182	329	1
## 183	330	2
## 184	331	1
## 185	333	1
## 186	335	2
## 187	336	1

```
## 188      337  1
## 189      338  1
## 190      339  2
## 191      340  2
## 192      341  3
## 193      342  3
## 194      344  1
## 195      347  1
## 196      349  1
## 197      353  1
## 198      354  1
## 199      355  1
## 200      358  1
## 201      360  1
## 202      365  1
## 203      369  1
## 204      384  1
## 205      385  1
## 206      388  1
## 207      392  1
## 208      393  1
## 209      394  2
## 210      404  1
## 211      407  1
## 212      409  1
## 213      412  1
## 214      417  1
## 215      458  1
## 216      466  1
## 217      468  1
## 218      491  1
## 219      518  1
## 220      529  1
## 221      564  1
## 222      603  1
```

```
heart %>%
  count(RestingBP)
```

```
##   RestingBP  n
## 1         0  1
## 2        80  1
## 3        92  1
## 4        94  2
## 5        95  6
## 6        96  1
## 7        98  1
## 8       100 15
## 9       101  1
## 10      102  3
## 11      104  3
## 12      105  9
## 13      106  3
## 14      108  7
## 15      110 58
```

## 16	112	14
## 17	113	1
## 18	114	2
## 19	115	19
## 20	116	2
## 21	117	1
## 22	118	10
## 23	120	132
## 24	122	12
## 25	123	2
## 26	124	12
## 27	125	29
## 28	126	7
## 29	127	1
## 30	128	18
## 31	129	1
## 32	130	118
## 33	131	4
## 34	132	17
## 35	133	6
## 36	134	11
## 37	135	20
## 38	136	13
## 39	137	5
## 40	138	17
## 41	139	5
## 42	140	107
## 43	141	3
## 44	142	11
## 45	143	2
## 46	144	8
## 47	145	18
## 48	146	4
## 49	148	2
## 50	150	55
## 51	152	7
## 52	154	3
## 53	155	8
## 54	156	2
## 55	158	4
## 56	160	50
## 57	164	1
## 58	165	2
## 59	170	14
## 60	172	2
## 61	174	1
## 62	178	3
## 63	180	12
## 64	185	1
## 65	190	2
## 66	192	1
## 67	200	4

```
heart %>%
  count(Age)
```

```
##      Age  n
## 1    28   1
## 2    29   3
## 3    30   1
## 4    31   2
## 5    32   5
## 6    33   2
## 7    34   7
## 8    35  11
## 9    36   6
## 10   37  11
## 11   38  16
## 12   39  15
## 13   40  13
## 14   41  24
## 15   42  18
## 16   43  24
## 17   44  19
## 18   45  18
## 19   46  24
## 20   47  19
## 21   48  31
## 22   49  21
## 23   50  25
## 24   51  35
## 25   52  36
## 26   53  33
## 27   54  51
## 28   55  41
## 29   56  38
## 30   57  38
## 31   58  42
## 32   59  35
## 33   60  32
## 34   61  31
## 35   62  35
## 36   63  30
## 37   64  22
## 38   65  21
## 39   66  13
## 40   67  15
## 41   68  10
## 42   69  13
## 43   70   7
## 44   71   5
## 45   72   4
## 46   73   1
## 47   74   7
## 48   75   3
## 49   76   2
## 50   77   2
```

```
heart %>%
  count(MaxHR)
```

```
##      MaxHR  n
## 1         60  1
## 2         63  1
## 3         67  1
## 4         69  1
## 5         70  1
## 6         71  1
## 7         72  2
## 8         73  1
## 9         77  1
## 10        78  1
## 11        80  2
## 12        82  3
## 13        83  1
## 14        84  3
## 15        86  4
## 16        87  1
## 17        88  2
## 18        90  3
## 19        91  1
## 20        92  6
## 21        93  2
## 22        94  4
## 23        95  2
## 24        96  7
## 25        97  3
## 26        98  9
## 27        99  7
## 28       100 14
## 29       102  4
## 30       103  4
## 31       104  2
## 32       105 11
## 33       106  5
## 34       107  1
## 35       108  8
## 36       109  5
## 37       110 23
## 38       111  5
## 39       112 13
## 40       113  5
## 41       114  6
## 42       115 16
## 43       116  9
## 44       117  6
## 45       118 12
## 46       119  5
## 47       120 36
## 48       121  5
## 49       122 20
## 50       123  7
```


## 51	124	9
## 52	125	21
## 53	126	12
## 54	127	8
## 55	128	14
## 56	129	4
## 57	130	33
## 58	131	7
## 59	132	11
## 60	133	5
## 61	134	6
## 62	135	15
## 63	136	6
## 64	137	7
## 65	138	14
## 66	139	6
## 67	140	41
## 68	141	6
## 69	142	14
## 70	143	10
## 71	144	13
## 72	145	14
## 73	146	6
## 74	147	5
## 75	148	11
## 76	149	6
## 77	150	43
## 78	151	5
## 79	152	11
## 80	153	5
## 81	154	12
## 82	155	14
## 83	156	10
## 84	157	7
## 85	158	8
## 86	159	5
## 87	160	25
## 88	161	7
## 89	162	13
## 90	163	10
## 91	164	4
## 92	165	11
## 93	166	5
## 94	167	2
## 95	168	8
## 96	169	6
## 97	170	20
## 98	171	4
## 99	172	10
## 100	173	7
## 101	174	7
## 102	175	10
## 103	176	2
## 104	177	1

```
## 105 178 6
## 106 179 6
## 107 180 10
## 108 181 2
## 109 182 6
## 110 184 4
## 111 185 4
## 112 186 2
## 113 187 1
## 114 188 2
## 115 190 2
## 116 192 1
## 117 194 1
## 118 195 1
## 119 202 1
```

```
heart %>%
  count(Oldpeak)
```

```
##      Oldpeak    n
## 1      -2.6     1
## 2      -2.0     1
## 3      -1.5     1
## 4      -1.1     1
## 5      -1.0     2
## 6      -0.9     1
## 7      -0.8     1
## 8      -0.7     1
## 9      -0.5     2
## 10     -0.1     2
## 11       0.0  368
## 12       0.1   14
## 13       0.2   22
## 14       0.3   11
## 15       0.4   11
## 16       0.5   19
## 17       0.6   14
## 18       0.7    7
## 19       0.8   16
## 20       0.9    4
## 21       1.0   86
## 22       1.1    7
## 23       1.2   26
## 24       1.3    7
## 25       1.4   18
## 26       1.5   53
## 27       1.6   16
## 28       1.7    6
## 29       1.8   17
## 30       1.9    7
## 31       2.0   76
## 32       2.1    2
## 33       2.2    5
## 34       2.3    2
## 35       2.4    4
```

```
## 36      2.5  16
## 37      2.6   7
## 38      2.8   7
## 39      2.9   1
## 40      3.0  28
## 41      3.1   1
## 42      3.2   2
## 43      3.4   3
## 44      3.5   2
## 45      3.6   4
## 46      3.7   1
## 47      3.8   1
## 48      4.0   8
## 49      4.2   2
## 50      4.4   1
## 51      5.0   1
## 52      5.6   1
## 53      6.2   1
```

```
heart %>%
  count(ST_Slope)
```

```
##   ST_Slope   n
## 1     Down  63
## 2     Flat 460
## 3       Up 395
```

```
heart %>%
  count(HeartDisease)
```

```
##   HeartDisease   n
## 1             0 410
## 2             1 508
```

```
heart <- heart %>%
  filter(Cholesterol != 0) %>%
  filter(RestingBP!=0) %>%
  mutate(Sex = as.factor(Sex)) %>%
  mutate(ChestPainType= as.factor(ChestPainType)) %>%
  mutate(RestingECG= as.factor(RestingECG)) %>%
  mutate(ExerciseAngina = as.factor(ExerciseAngina)) %>%
  mutate(ST_Slope = as.factor(ST_Slope)) %>%
  mutate(HeartDisease = as.factor(HeartDisease))
```

```
#set.seed(1)
#x <- model.matrix(HeartDisease~., heart)[,-1]
#y <- heart$HeartDisease

#sumfit <- sum(y ~ ., data = heart, kernel = "linear", cost = 10, scale = FALSE)
#summary(sumfit)
```

The code below creates a training and test data set.

```
set.seed(1)
training <- sample(dim(heart)[1], 522)
train <- heart[training, ]
test <- heart[-training, ]
```

The code below fits a support vector classifier.

```
svm.fit <- svm(HeartDisease ~., data = train, kernel = "linear", cost=0.01)
summary(svm.fit)
```

```
##
## Call:
## svm(formula = HeartDisease ~ ., data = train, kernel = "linear",
##      cost = 0.01)
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##      cost:   0.01
##
## Number of Support Vectors: 295
##
## ( 148 147 )
##
##
## Number of Classes: 2
##
## Levels:
## 0 1
```

There are 295 Support Vectors, 148 in Class 0 and 147 in Class 1.

```
svm.pred <- predict(svm.fit, test)
wrong <- 0
for (i in 1:length(svm.pred)){
  if (svm.pred[i]!=test$HeartDisease[i]){
    wrong = wrong + 1
  }
}
print(wrong/length(test$HeartDisease))
```

```
## [1] 0.1428571
```

The test error rate is 14.3%.

The code below uses CV to find the optimal cost for an svm with a linear kernel.

```
set.seed(1)
tune.out = tune(svm, HeartDisease ~ ., data = train, kernel = "linear", ranges = list(cost = 10^seq(-2,
summary(tune.out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
## 5.623413
##
## - best performance: 0.1398403
##
```

```
## - Detailed performance results:
##      cost      error dispersion
## 1  0.01000000 0.1532656 0.04594129
## 2  0.01778279 0.1532656 0.04021823
## 3  0.03162278 0.1571118 0.04398593
## 4  0.05623413 0.1417634 0.03134359
## 5  0.10000000 0.1474964 0.02842013
## 6  0.17782794 0.1417271 0.02546022
## 7  0.31622777 0.1436502 0.02558931
## 8  0.56234133 0.1474601 0.03084972
## 9  1.00000000 0.1493832 0.03187442
## 10 1.77827941 0.1417634 0.03134359
## 11 3.16227766 0.1436865 0.03011101
## 12 5.62341325 0.1398403 0.03110896
## 13 10.00000000 0.1398403 0.03110896
```

The optimal cost is 5.623.

The code below calculates the test error using the optimal cost.

```
svm.fit = svm(HeartDisease ~ ., kernel = "linear", data = train, cost = tune.out$best.parameters$cost)

svm.pred <- predict(svm.fit, test)

wrong <- 0
for (i in 1:length(svm.pred)){
  if (svm.pred[i] != test$HeartDisease[i]){
    wrong = wrong + 1
  }
}
print(wrong/length(test$HeartDisease))

## [1] 0.1205357
```

The error rate is 12.05%.

END OF LINEAR

BEGINNING OF RADIAL

The code below fits a support vector classifier.

```
svm.fitradial <- svm(HeartDisease ~ ., data = train, kernel = "radial")
summary(svm.fitradial)

##
## Call:
## svm(formula = HeartDisease ~ ., data = train, kernel = "radial")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##      cost:   1
##
## Number of Support Vectors: 233
##
## ( 119 114 )
```

```
##
##
## Number of Classes: 2
##
## Levels:
## 0 1
```

There are 233 Support Vectors, with 119 in 0 and 114 in 1.

```
svm.predradial <- predict(svm.fitradial, test)

wrong <- 0
for (i in 1:length(svm.predradial)){
  if (svm.predradial[i]!=test$HeartDisease[i]){
    wrong = wrong + 1
  }
}
print(wrong/length(test$HeartDisease))
```

```
## [1] 0.1205357
```

The test error is 12.05%.

```
set.seed(1)
tune.out2 = tune(svm, HeartDisease ~ ., data = train, kernel = "radial", ranges = list(cost = 10^seq(-2
summary(tune.out2)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##      cost
## 5.623413
##
## - best performance: 0.1399129
##
## - Detailed performance results:
##      cost      error dispersion
## 1  0.01000000 0.5171263 0.07196263
## 2  0.01778279 0.2298621 0.05191857
## 3  0.03162278 0.1800798 0.04793246
## 4  0.05623413 0.1628084 0.04793032
## 5  0.10000000 0.1513788 0.04653122
## 6  0.17782794 0.1513788 0.04740609
## 7  0.31622777 0.1551887 0.04380499
## 8  0.56234133 0.1475327 0.03145408
## 9  1.00000000 0.1437591 0.03657071
## 10 1.77827941 0.1418723 0.03554973
## 11 3.16227766 0.1456821 0.04276003
## 12 5.62341325 0.1399129 0.04156757
## 13 10.00000000 0.1456821 0.04276003
```

The optimal cost is 5.623. The code below finds the test error of a radial kernel SVM using the optimal cost.

```

svm.radial3 = svm(HeartDisease ~ ., data = train, kernel = "radial", cost = tune.out2$best.parameters$cost)

svm.predradial <- predict(svm.radial3, test)

wrong <- 0
for (i in 1:length(svm.predradial)){
  if (svm.predradial[i]!=test$HeartDisease[i]){
    wrong = wrong + 1
  }
}
print(wrong/length(test$HeartDisease))

```

```
## [1] 0.1205357
```

The test error is 12.05%.

END OF RADIAL

BEGINNING OF POLYNOMIAL

```

svm.fitpoly <- svm(HeartDisease ~ ., data = train, kernel = "poly", degree = 2)
summary(svm.fitpoly)

```

```

##
## Call:
## svm(formula = HeartDisease ~ ., data = train, kernel = "poly", degree = 2)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: polynomial
##       cost:  1
##   degree:  2
##   coef.0:  0
##
## Number of Support Vectors:  279
##
##   ( 139 140 )
##
##
## Number of Classes:  2
##
## Levels:
##   0 1

```

There are 279 Support Vectors, 139 0 and 140 1.

```

svm.predpoly <- predict(svm.fitpoly, test)

wrong <- 0
for (i in 1:length(svm.predpoly)){
  if (svm.predpoly[i]!=test$HeartDisease[i]){
    wrong = wrong + 1
  }
}
print(wrong/length(test$HeartDisease))

```

```
## [1] 0.1473214
```

The test error is 14.7%.

The code below find teh optimal cost

```
set.seed(1)
tune.out = tune(svm, HeartDisease ~ ., data = train, kernel = "poly", degree = 2, ranges = list(cost = 10.00000000))
summary(tune.out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##      cost
## 1.778279
##
## - best performance: 0.1456821
##
## - Detailed performance results:
##      cost      error dispersion
## 1  0.01000000 0.5171263 0.07196263
## 2  0.01778279 0.5171263 0.07196263
## 3  0.03162278 0.4518505 0.07328965
## 4  0.05623413 0.2813498 0.06687009
## 5  0.10000000 0.2183962 0.03164476
## 6  0.17782794 0.1781930 0.04068122
## 7  0.31622777 0.1628447 0.03866025
## 8  0.56234133 0.1532656 0.03949555
## 9  1.00000000 0.1494920 0.03938606
## 10 1.77827941 0.1456821 0.03872582
## 11 3.16227766 0.1533745 0.04266201
## 12 5.62341325 0.1572206 0.04881518
## 13 10.00000000 0.1725327 0.05379036
```

The optimal cost is 1.778.

The code below calculates the test error using the optimal cost.

```
svm.poly3 = svm(HeartDisease ~ ., data = train, kernel = "poly", degree = 2, cost = tune.out$best.parameters$cost)

test.pred3 = predict(svm.poly3, test)

wrong <- 0
for (i in 1:length(test.pred3)){
  if (test.pred3[i] != test$HeartDisease[i]){
    wrong = wrong + 1
  }
}
print(wrong/length(test$HeartDisease))
```

```
## [1] 0.1294643
```

The test error is 12.95%.

CONCLUSION

Linear Kernel w/Optimal Cost: 12.05%

Radial Kernel w/Optimal Cost: 12.05%

Polynomial Kernel w/ Optimal Cost: 12.95%
